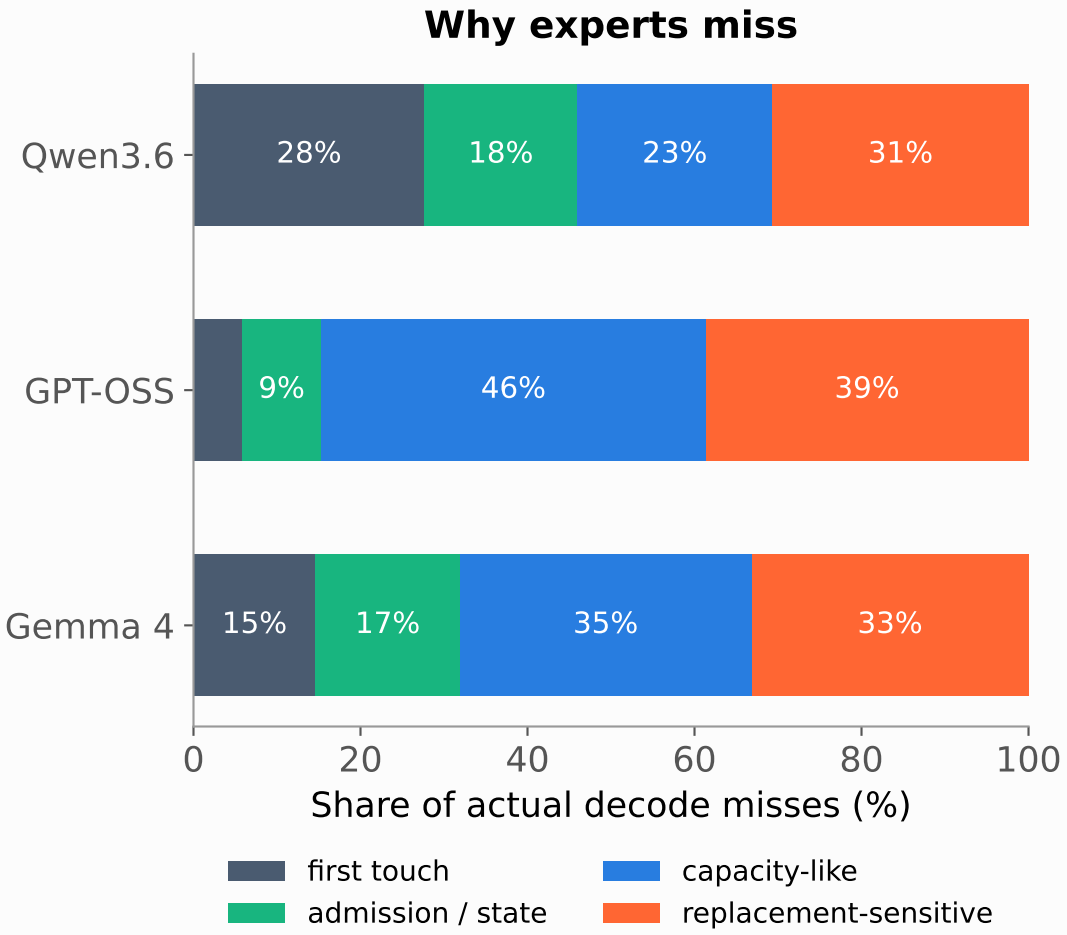
Why experts miss
Qwen3.6
GPT-OSS
Gemma 4
28%
18%
23%
31%
9%
46%
39%
15%
17%
35%
33%
0
20
40
60
80
100
Share of actual decode misses (%)
first touch
admission / state
capacity-like
replacement-sensitive

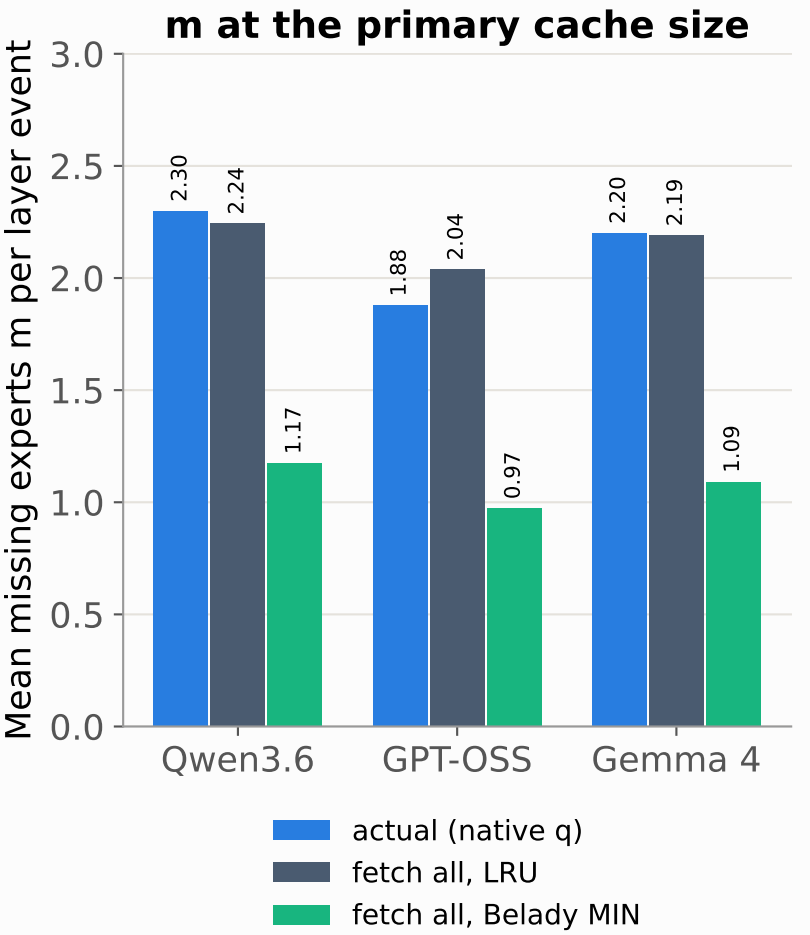
m at the primary cache size
Mean missing experts m per layer event
3.0
2.5
2.0
1.5
1.0
0.5
0.0
2.30
2.24
1.17
1.88
2.04
0.97
2.20
2.19
1.09
Qwen3.6
GPT-OSS
Gemma 4
actual (native q)
fetch all, LRU
fetch all, Belady MIN