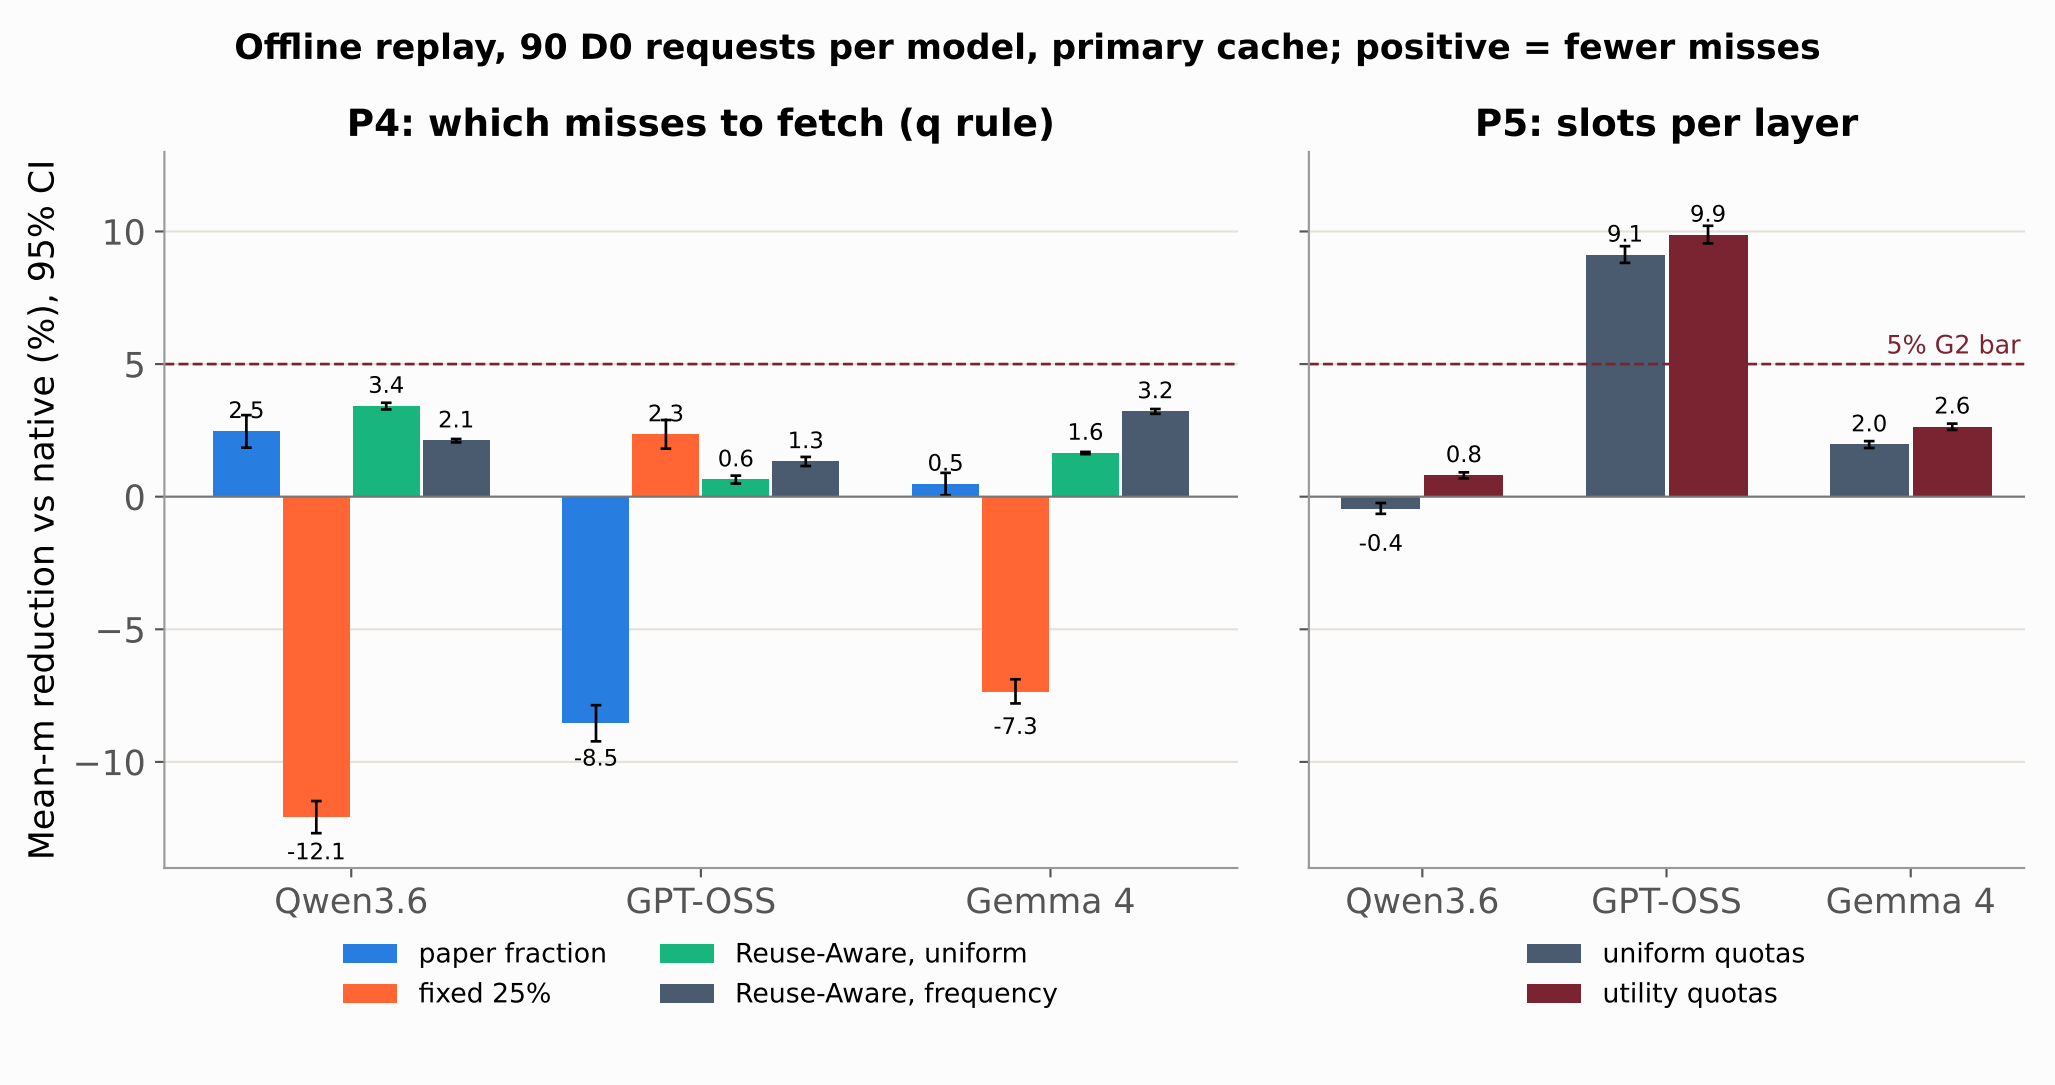
Offline replay, 90 D0 requests per model, primary cache; positive = fewer misses
P4: which misses to fetch (q rule)
P5: slots per layer
Mean-m reduction vs native (%), 95% CI
10
5
0
−5
−10
2.5
-12.1
3.4
2.1
-8.5
2.3
0.6
1.3
0.5
-7.3
1.6
3.2
-0.4
0.8
9.1
9.9
2.0
2.6
5% G2 bar
Qwen3.6
GPT-OSS
Gemma 4
paper fraction
fixed 25%
Reuse-Aware, uniform
Reuse-Aware, frequency
uniform quotas
utility quotas